

Wine Quality Classification

Dataset: Wine-quality-red

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Location of dataset
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/
    winequality-red.csv"
# Reading dataset to pandas dataframe
data = pd.read_csv(url, sep=";")
```

```
[3]: data.head()
```

```
[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5

1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[4]: data.describe()
```

```
[4]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806
std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

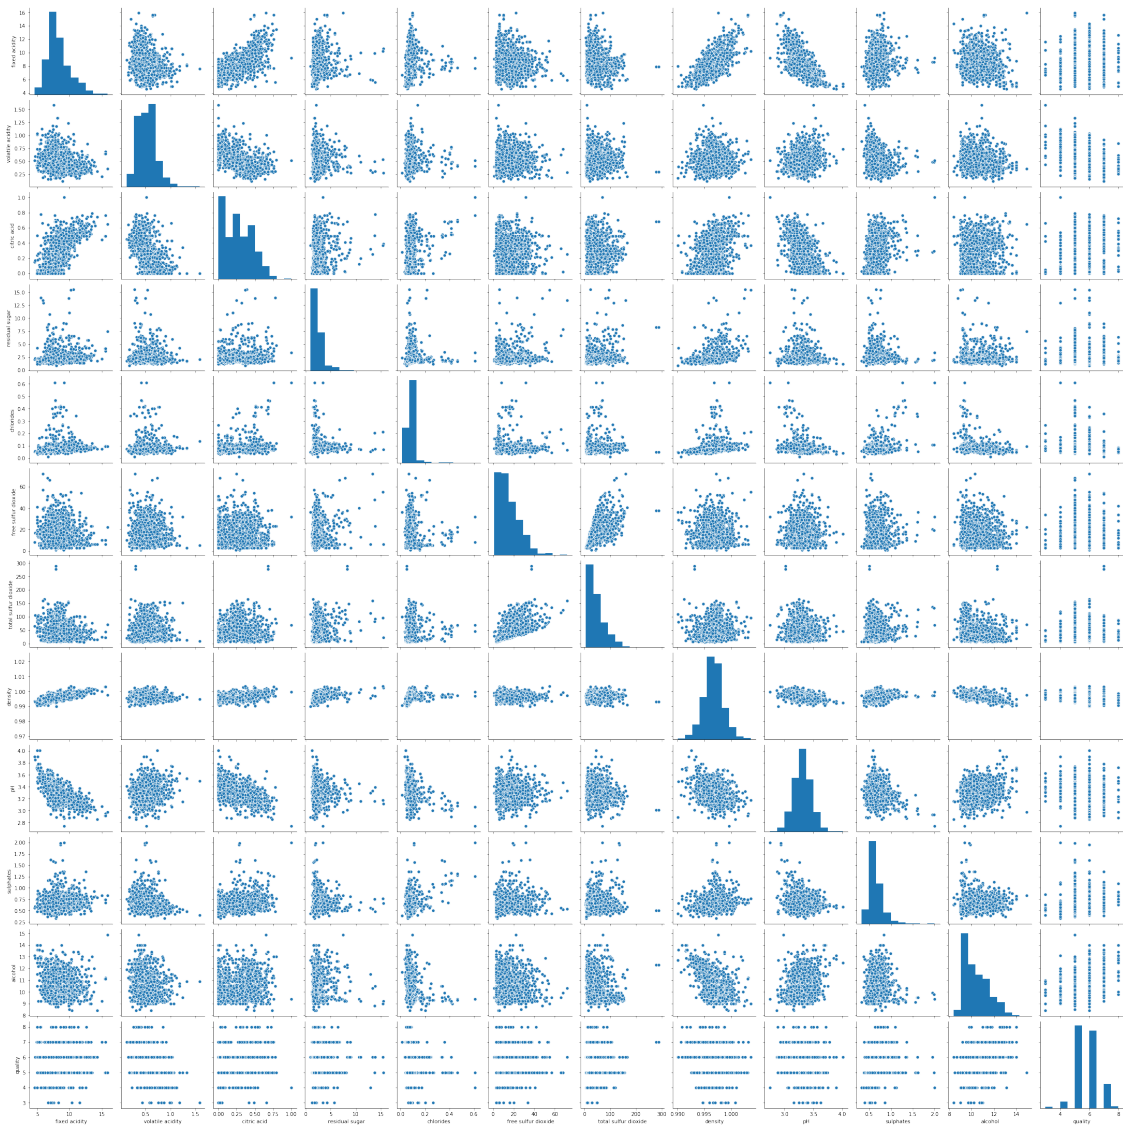
```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
```


```
3  residual sugar      1599 non-null  float64
4  chlorides           1599 non-null  float64
5  free sulfur dioxide  1599 non-null  float64
6  total sulfur dioxide 1599 non-null  float64
7  density             1599 non-null  float64
8  pH                  1599 non-null  float64
9  sulphates           1599 non-null  float64
10 alcohol             1599 non-null  float64
11 quality             1599 non-null  int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
[6]: # Visualizing the variables
plt.figure(figsize=(25,25))
sns.pairplot(data)
plt.show()
```

<Figure size 1800x1800 with 0 Axes>



```
[7]: # Correlation among variables
plt.figure(figsize=(25,20))
sns.heatmap(data.corr(),annot=True)
plt.show()
```




```
[8]: # Checking if any null values are there in dataset
data.isnull().sum()
```

```
[8]: fixed acidity      0
      volatile acidity  0
      citric acid       0
      residual sugar    0
      chlorides         0
      free sulfur dioxide 0
      total sulfur dioxide 0
      density           0
      pH                0
      sulphates         0
      alcohol           0
      quality           0
      dtype: int64
```

```
[9]: # Checking the quality of wine
data['quality'].value_counts()
```

```
[9]: 5    681
      6    638
      7    199
      4     53
      8     18
      3     10
      Name: quality, dtype: int64
```

```
[10]: # Plotting the quality of wine
plt.figure(figsize=(10,5))
sns.countplot(data['quality'])
plt.show()
```




```
[11]: # Classifying wine quality as good and bad (anything below 7 as bad, and from 7
      ↪ is good)
bins = (2, 6.99, 8)
group_names = ['bad', 'good']
data['quality'] = pd.cut(data['quality'], bins=bins, labels=group_names)
data['quality'].unique()
```

```
[11]: [bad, good]
Categories (2, object): [bad < good]
```

```
[12]: # Labeling bad wine quality as 0 and good wine quality as 1
label_quality = LabelEncoder()
data['quality'] = label_quality.fit_transform(data['quality'])
data['quality'].value_counts()
```

```
[12]: 0    1382
      1     217
      Name: quality, dtype: int64
```

```
[13]: # Plotting the target variables
plt.figure(figsize=(10,5))
sns.countplot(data['quality'])
plt.show()
```



```
[14]: # Assign data from first 11 columns to X variable
X = data.drop('quality', axis=1)
# Assign data from last 12th column to y variable
y = data['quality']
```

```
[15]: # Splitting the dataset into (3:1) train-test ratio
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.75,
→test_size = 0.25)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1199, 11)
(400, 11)
(1199,)
(400,)
```

```
[16]: # Feature scaling (by standardization)
```

```
scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
print(X_train)
print(X_test)
```

```
[[-0.07563803 -0.97403017  0.3002543  ... -0.25332521  0.965028
  0.55331102]
 [-0.88744108  0.43787425 -0.36651643 ...  1.52214504 -0.46467959
 -1.04438534]
 [-0.30758176 -0.15946993  0.14638413 ...  0.00970742 -0.58900199
 -0.01058181]
 ...
 [-1.6412582  -0.21377394 -1.23844737 ...  1.25911241  1.0271892
  3.37277519]
 [ 2.41775706 -1.35415828  1.37734547 ... -1.30545573  4.38389397
 -0.57447464]
 [-0.30758176 -0.7568141  0.7618648  ...  0.53577268  0.03261
  1.11720386]]
[[-1.06139888 -0.7568141  -0.21264626 ...  0.86456347 -0.65116319
  1.58711455]
 [-0.42355362 -0.64820607  0.3002543  ... -0.05605074 -0.46467959
  0.7412753 ]
 [-0.53952549  0.57363429 -0.87941698 ...  1.78517767  0.2190936
  0.36534675]
 ...
 [ 0.27227756 -0.64820607  1.89024602 ...  1.52214504  0.03261
  0.08340033]
 [ 0.3302635  -0.26807796  1.32605541 ... -0.64787416 -0.40251839
  1.58711455]
 [ 0.3302635  -0.15946993 -0.31522637 ...  0.14122373  1.0893504
  0.64729316]]
```

```
[32]: # Multilayer Perceptron Classifier
```

```
mlp = MLPClassifier(hidden_layer_sizes=(100,1000,100), max_iter=10000)
mlp.fit(X_train, y_train.values.ravel())
```

```
[32]: MLPClassifier(hidden_layer_sizes=(100, 1000, 100), max_iter=10000)
```

```
[33]: predictions = mlp.predict(X_test)
predictions
```

```
[33]: array([0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```



```

0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0,
1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0, 0, 1, 0])

```

```
[34]: print(confusion_matrix(y_test, predictions))
```

```

[[325  23]
 [ 17  35]]

```

```
[35]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	348
1	0.60	0.67	0.64	52
accuracy			0.90	400
macro avg	0.78	0.80	0.79	400
weighted avg	0.91	0.90	0.90	400

```
[36]: print("Accuracy: " + str(accuracy_score(y_test, predictions)*100) + "%")
```

```
Accuracy: 90.0%
```